# Algorithms For Ricochet Robots

Student: Anders Rydbirk, s113725

Advisors: Inge Li Gørtz and Philip Bille

Abstract

The topic of this project is algorithms for the game *Ricochet Robots*. Optimizing the cost of a single entity is the most common approach in path finding problems, but the game has several characteristics that challenges this basic assumption. In this project two existing and three new algorithms are presented, analysed, implemented, and performance tested. The algorithms all have different approaches to the key problems of limiting the search tree and guaranteeing an optimal solution. Two of the new algorithms presented construct a graph of the board with dynamic progamming and performs a BFS in incrementing steps while the last solves the game with a worst case running time only depending on the size of the board. The new algorithms have all shown great results during the performance test at a minor tradeoff in guaranteeing an optimal solution. The findings in this project also reveals possible improvements in the algorithms allowing them to increase the percentage of correct solutions while still being capable of competing with existing solutions.

Contents

1 Introduction

Optimizing the cost of a single entity is the most common approach in path finding problems. However, the board game *Ricochet Robots* presents a rather unusual case where multiple entities coorperate for a single entity to reach the goal at minimal cost. Usually, the map is considered static in path finding problems, however, in *Ricochet Robots* the map is altered for each move of an entity, resulting in a new problem. With multiple possible moves for each map, the search space increases rapidly.

The goal of this project is to survey and design state-of-the-art algorithms within graph search and dynamic programming. I will survey existing solutions for *Ricochet Robots* and, based on the experience, design and implement new state-of-the-art algorithms. I will be implementing prototypes of each algorithm and analyse and evaluate each algorithm from both a theoretical and practival perspective.

1.1 Ricochet Robots

Ricochet Robots is played on a board consisting of $16 \cdot 16$ tiles with obstacles scattered on the board as seen in figure 1. The board contains 17 goals. Four robots in different colors are randomly placed on the board at the beginning of the game. The robots cannot be positioned on the same field or on a goal field. Each goal is colored in relation to one of the four robots except for one of the goals which is reachable by all colors. The game is played by solving the individual goal sequentially. The objective of the game is to identify the minimum number of moves on order for the corresponding robot to land on the given goal. All robot can be moved and all robots can be used as obstacles for the other robots.

A round starts with picking a goal to reach. All players can declare they are able to reach the goal in a specific number of moves. When the first player declares a number, all the other players have one minute to find a lower (and more efficient) number of moves. If so, the player calling a lower number of moves demonstrates his/her solution and wins the round, and if not, the declaring player wins the round. The game continues until all 17 goals have been played. The game is now over and the aggregated winner is found.

Moving the robots is one of the characteristic parts of the game. A robot moves like a rook in chess, but it moves in a direction until it hits an obstacle or another robot. This counts as one move. Thus, a robot cannot stop on any of the intermediate fields. All robots can be moved and used as obstacles for the other robots. An example is given in figure 1. The robots can move in four directions. As a result, for each move, $|directions| \cdot |robots| = 16$ new possible moves can be made. This is also refered to as the branching factor which is also illustrated in figure 2. However, as the robots move they are adjacent to at

Figure 1: An example of *Ricochet Robots* where the green robot has to reach the goal at (7,2). The minimum number of moves is seven, and multiple solutions exist. The number on each arrow indicates the order of moves. Thus, the yellow robot moves first so that the green robot can bounce off it on the fourth move.

least one obstacle, and the branching factor appromixates 12. The number of possible combinations for a given game is given by 12^k where k is the required number of moves to reach the goal. Given the exponential growth in possible combinations, the game presents three major challenges for both players and computers:

1. Finding a valid solution
2. Guaranteeing an optimal solution
3. Solving the game in reasonable time

The game has been proven NP-Hard [2] which underlines the complexity of the game is an exponential function of the required number of moves. Precomputing each and every robot combination for each board is infeasible given the high number of robot combinations and the complexity of each game. Thus, tradeoffs between solving the game optimally and solving the game in reasonable time are the core of developing algorithms for this game.

1.2 Baseline

The following definitions are general for all algorithms.

Figure 2: The branching factor illustrated. In this case, the branching factor is 12 while each circle is a game state. Therefore, for each move conducted, 12 new possible moves can be made.

Let b be the branching factor, and let k be the number of moves to reach the goal. The number of possible moves in each game is given by b^k.

Let n be the dimensions of the board where *n = 16*. Let F be the set of fields where *F[i,j]* refers to the field at the $i'th$ row and the $j'th$ column starting from the top-left corner. Each field contains information about obstacles in each direction. Let R be the set of robots, let GR be the Goal Robot to reach the goal and let OR be the set of Obstacle Robots $\in R - \{GR\}$.

A robot state refers to the position of the robot and the required moves to reach the state, while a game state refers to the complete set of robot states for a current configuration.

Let b be the branching factor describing the number of possible moves for each game state and let k be the minimum number of moves required for solving a given game.

The set of directional indicators D refers to the possible directions on the game-board. Thus, $D = \{North, East, South, West\}$.

Each algorithm is given the board of fields F, the set of robots R and the goal to reach G as input.

1.3 Previous Solutions

A few solutions exist for *Ricochet Robots* and the two most common are described below.

1.3.1 Naive Algorithm

The naive algorithm searches the entire search tree in an incremental order and guarantees an optimal solution. The board representation is a two dimensional

array with an additional attribute for each field indicating if a robot stands on the given field. A first-in, first-out queue is used for keeping track of the to-be processed game states. Since this solution would never be feasable for solutions requiring more than a couple moves, a hash table is introduced to keep track of already searched game states. Therefore, duplicate game states will be pruned from the search tree.

To move a robot, all fields in a given direction is processed until an obstacle is met. It uses $O(n)$ time. This is done for each robot for each direction for each game state. Processing each game state includes the branching factor and uses $O(b \cdot n)$ time.

Solving a given configuration uses $O(b^k \cdot b \cdot n) = O(b^{k+1} \cdot n)$ time. The algorithm uses $O(n^2)$ space for the additional attribute for each field, $O(b^k)$ space for the queue and $O\left(\sum_{i=1}^{k} b^i \right) \sim O(b^k)$ for the hash table. The worst case space impact of the algorithm is $O(b^k + n^2)$.

Each robot is adjacent to at least one obstacle after the initial move. Therefore the branching factor is given by $b = |R| \cdot (|D| - 1) = 12$.

1.3.2 Iterative Deepening Depth First Search (IDDFS)

The IDDFS algorithm is claimed to be fastest solver and guarantees an optimal solution [1]. The algorithm uses the same board representation as the naive algorithm including the additional attribute for flagging robot locations. In the following, h will refer to the height of the remaining search for the iteration. Thus, $h = MAX - depth$, where MAX is the incrementing search limit while *depth* is the distance to the root in the search tree. The IDDFS includes two additional pruning techniques:

1. A hash table stores a combination of game state s and h. If s has been reached before by the same height, it is pruned from the search.
2. A two-dimensional array stores the minimum number of moves from each field to goal. In the latter, *min[i,j]* refers to the minimum number of moves to goal from position *(i,j)*. It is assumed that the robots can change direction without bouncing off an obstacle as seen in figure 3. If $min[GR] > h$ the search is pruned. If $h = min[GR]$ only the GR is processed further in the search tree.

The minimum moves for each field is computed with the goal as the root. The number of moves for the root is set to 0. Each direction is processed incrementing the number of moves with 1. For each direction, all fields is processed until an obstacle or a field with a lower minimum number of moves is met. Due to this incremental movement, every field is only queued once but can be visited several times during the computation. When processing a single field, only two

3	3	3	2	2	2	2	2	2	2	2	2	3	3	3	3
3	3	0	1	1	1	1	1	1	1	1	1	3	3	3	4
2	2	1	2	2	2	2	2	2	2	2	2	2	2	2	2
2	2	1	2	2	2	2	2	2	2	2	2	3	3	3	3
2	2	1	2	2	2	2	2	2	2	2	2	2	2	2	3
2	2	1	2	2	2	3	2	2	3	3	2	3	3	3	3
3	3	1	2	2	2	2	2	2	2	2	2	2	2	2	2
2	2	1	2	2	2	2	■	■	3	3	2	3	3	3	3
2	2	1	2	2	2	2	■	■	3	3	2	3	3	3	3
2	2	1	2	3	2	3	3	3	3	3	2	3	3	3	3
2	2	1	2	2	2	2	2	2	2	2	2	2	2	2	2
3	2	1	2	2	2	2	2	2	2	3	2	3	3	3	3
2	2	1	2	2	2	2	3	3	3	3	2	3	3	3	3
2	2	1	2	2	2	2	2	2	2	2	2	2	2	2	2
4	3	3	3	3	2	3	3	3	3	3	2	3	4	3	3
3	3	3	3	3	2	3	3	3	3	3	2	3	3	3	3

Figure 3: The minimum number of moves for each field to reach the goal at (2,3) colored in grey. All moves assumes optimal placement of other robots.

directions is relevant since the field was discovered from one of the other two directions and therefore these directions are guaranteed to be more optimal. When processing a single field, only $O(n)$ fields are visited. Since all fields are queued exactly once, the minimum moves are computed in $O(n^2 \cdot n) = O(n^3)$ time and uses $O(n^2)$ space.

Finding the optimal solution uses $O(b^k \cdot n + n^3)$ time. However, the space impact of the search itself is only $O(k)$ since the algorithm is recursive. The space analysis of the hash table is the same as in the naive algorithm. The worst case space impact is $O(b^k + k + n^2) = O(b^k + n^2)$.

2   Solutions

In the following, I will describe my solutions to the problem presented by *Ricochet Robots* and analyse each for running time and space impact.

2.1   Stateless

The stateless algorithm is an attempt to solve the game without keeping track of each game state. The obvious advantage of this is to eliminate impact of the branching factor on the running time. The algorithm does not guarantee an optimal solution. Instead, it stores information about every time a robot

interacts with a field. The interaction can be split in two - when the robot moves over the field, or when it lands on the field. In the latter example, landing on a field is denoted as a *final state* while moving over a field is denoted as an *intermediate state*. The algorithm assumes that every field can only have one final state for each robot, and only one intermediate state for each combination of direction and color.

Final states are queued in a min-heap queue ordered by the number of moves. These states represent only a single state of a robot, thus for each element in the queue, only one robot is processed. The robot is moved in all directions and a new state are created for each visited field if no redundant state already exist for the given field. The algorithm terminates when there are no more robot states to process, whereafter the best result is returned.

2.1.1 The Stateless Algorithm

Let *S[i,j]* refer to the element storing states for *F[i,j]*. *S[i,j]* is instantiated at the first visit to *F[i,j]*. Each $s \in S$ stores final states in a hash table using the color of the robot as key. In addition, *s* stores all intermediate states in a hash table using the direction of the robot move as the key. Finally, *s* stores if a robot starts on the field.

2.1.1.1 Moving And States

Moving a robot determines the state for each field processed. The state depends upon the next field to process if it is possible to move further. Four cases are taken into consideration when moving robot *r*, and all four cases are evaluated for each step until an obstacle is met. In the following, *a* will refer to the adjacent field in direction *d*:

1. *F[i,j]* cannot move in direction *d*. A final state *f* is added to *S[i,j]*. If a final state did not already exist, *f* is added to the queue for future processing.
2. *F[i,j]* can move in direction *d*:
    a. No final robot states exist on *a*. Thus, an intermediate state is added to *S[i,j]* and *a* is processed next.
    b. A final robot state exist on *a* and its color does not match the color of robot *r*. The final state *f* is set to depend upon the lowest final state in *S[a]*, the move counts are aggregated and *f* is added to *S[i,j]*. If a final state did not already exist, *f* is added to the queue while *a* is processed next.
    c. A robot starts on *a*. Case *1.b* is applied, with the addition that all the further processing of fields in direction *d* for the current *r* will be added one to the number of required moves from *a* and forward. This refers to the required movement of the starting robot on *a*.

9

Figure 4: The initial move for the blue robot. All grey fields are final states while all fields denoted with an appostrophe are intermediate states. The field of the red robot is also an intermediate state with a value of 2.

An initial move is illustrated in figure 4 displaying several of the cases.

Whenever a new final state f is added to the queue, adjacent fields are searched for intermediate states on collision with f. Let a be the adjacent field in direction d and let d' be the opposite direction of d. All intermediate states $i \in I$ where $i.color \neq f.color \wedge i.direction = d'$ are removed from $S[a]$ and converted to final states that depend upon f. If no final state already exists for the given color in $S[a]$, the new final states are added to the queue.

2.1.2 Analysis

The algorithm uses a min-heap queue and insertion and removal uses $O(log(m))$ time where m refers to the number of elements in the queue. Since each field is only processed once for every robot, $m = O(n^2 \cdot |R|)$. Thus, insertion and removal in the queue uses $O(log(|R| \cdot n^2))$ time. Moving a robot uses $O(2n)$ time like the naive algorithm. Inserting, lookup, and removal from the hash tables in S uses $O(1)$ time. Processing each robot state uses $O(2n \cdot 2log(|R| \cdot n^2))$ time and the total worst case uses $O(n^2 \cdot 2n \cdot 2log(|R| \cdot n^2)) = O(n^3 \cdot log(n^2))$ time.

In the worst case, each $s \in S$ stores $O(|R|)$ final states and $O(|D| \cdot |R|)$ intermediate states. The queue contains $O(|R| \cdot n^2)$ elements. Therefore, the worst case space impact is $O(n^2 \cdot (1 + |R| + |R| \cdot |D|)) = O(n^2)$.

2.2 Graph v1

In the following the graph based algorithm will be presented. The algorithm assumes that all moves by the *OR* can be applied before the *GR* moves towards the goal. The *GR* can depend on *OR* states but *OR* can only depend on the starting state of *GR*. An example of the implications of this assumption is shown in figure 5.


Figure 5: Case A displays the optimal solution for the given game board while case B is the optimal solution given by the graph algorithm hence case A requires the red robot to move after the initial move of the blue robot.

The algorithm will update the graph with *OR* states and perform a BFS with the *GR* as the source vertex and stores the result. Like the naive algorithm, each *OR* are now moved to new positions and new game states are queued in a first-in, first-out queue. As each level in the search tree is discovered in an incremental order, the algorithm guaranties an optimal solution under the given assumption above. As with the naive algorithm, a hash table is used to prune the search tree.

2.2.1 The Graph v1 Algorithm

Let *G* be the graph representing the board where *G.V* is the set of vertices in the graph where *G.V[i,j]* refers to the vertex at the same position as the corresponding field in *F[i,j]*. Let *G.E* be all directional edges in *G* and let *G.Adj[v]* refer to all edges where the source is the vertex *v*. *G* will be represented in a two-dimensional array with directional edges stored in an adjacency list for each vertex. Constructing the graph is done in two passes with dynamic programming. Let $v \in G.V$ and let $E = G.Adj[v]$. Each edge $e \in E$ has a pointer to the next vertex, later referred to as the child of *e*. Each *e* is also denoted a directional indicator $d \in D$. While *v* has at most one edge for each direction, then $|E| \leq |D|$ is given.

2.2.1.1   Construct Graph

The graph is constructed in two passes, where the first pass is conducted in a top-down left-to-right approach. The first pass processes edges in the directions $d \in \{West, North\}$. Each *G.V[i,j]* is added an edge in direction *d* if *F[i,j]* has no obstacle in direction *d*. In the case of $d \in \{West, North\}$, the respective vertices *G.V[i,j-1]* and *G.V[i-1,j]* will be investigated. As an example, in case *F[i,j]* can move in direction *d = North* and *G.V[i-1,j]* has an edge in direction *d*, the child of that edge will also be the child of the edge in direction *d* for vertex *G.V[i,j]*. If *G.V[i-1,j]* has no edge in direction *d*, *G.V[i-1,j]* will be the child of the edge at direction *d* for vertex *G.V[i,j]*. The second pass processes edges where $d \in \{ East, South \}$ in a bottom-up right-to-left approach.

2.2.1.2   Adapting The Graph

The graph has to be adapted to the current *OR*'s positions before the BFS can be applied for the *GR*. Let *(i,j)* be the position where a robot has to be placed. Let *d* be a direction from *(i,j)* and let *V* be the set of vertices in direction *d* which need to be adjusted to the new board configuration. All vertices in *V* need to update the edges in the opposite direction of *d*, in the following referred to as d'.

Let v' be the vertex immediately adjacent to *G.V[i,j]* in direction *d*. Let *E* be the set of $e \in G.Adj[V - \{v'\}]$ where $e.d = d'$. The child property of all edges in *E* has to be set to v' while the edge of v' in direction d' has to be removed.

To find the set *V*, the most distant vertex is found and then all vertices in direction d' is processed until *G.V[i,j]* is met. Several cases exist when finding the distant vertex which is illustrated in figure 6 and listed below.

1. *F[i,j]* cannot move in direction *d*. No edges has to be updated.
2. *F[i,j]* can move in direction *d*:

   a. *G.V[i,j]* has no edge in direction *d*. Thus, another robot stands on the adjacent field in direction *d*.
   b. *G.V[i,j]* has an edge *e* in direction *d*. The field for *e.child* cannot move in direction *d*. Thus, *e.child* is the most distant vertex from *G.V[i,j]* in direction *d*.
   c. *G.V[i,j]* has an edge *e* in direction *d*. The field for *e.child*, *f*, can move in direction *d*. Thus, a robot is present on the adjacent field to *f* in direction *d*.

Given the position *(i,j)*, the fields *F* and the graph *G*, all directions are processed and the graph is updated with the new robot position.

Placing and removing a robot from the board involves many of the same features:

1. Identify relevant directions to investigate.

Figure 6: Each case when updating the graph in direction d. The grey field is the most distant vertex in each case.

2. Identify relevant vertices to update.
3. Add or update edges for all affected vertices.

This functionality will not be described further in this section.

2.2.1.3   Solver

In the graph-based solver the search tree is expanded one level at a time. A first-in, first-out queue is used to store awaiting game states. In this case, a game state refers to a list of *OR* states and a property indicating the total number of moves for the given game state. Like the naive algorithm, a hash table is introduced to track already processed game states and prunes redundant states.

The solver keeps track of the best known solution so far in *B*. *B.s* is the game state, *B.bfs* is the result of the BFS and *B.best* is the combined number of moves for the given solution. A heuristic function is introduced as the termination state:

> B is the optimal solution if $B.best < (s.moves + 2)$, where *s* is the curent investigated game state. Since the number of moves is investigated in an incremental order, and since *s* requires at least one additional move for the goal robot to reach the goal, *s* will never be a better solution than *B*, and therefore the solver terminates when this condition is fulfilled.

The solver instantiates the hash table and constructs the graph. Then the initial game state is queued and a loop is performed until a solution is found. For each game state, the heuristic is applied. If the termination state is not fulfilled, the graph is updated with the the positions of the *OR* for the given game state. The BFS is applied and the result of the goal field is found. If it was reached by the *GR*, the combined result is compared to *B* and updated if better. Then, the graph is adapted to the position of the *GR*, and all the *OR* are moved in all

directions. The new game states are validated against the hash table and added to the queue if they are not redundant. Finally, the robots are all removed from the graph and the next game state is processed.

2.2.2 Analysis

Processing a single field in the construction of the graph uses $O(1)$ time, and each pass uses $O(n^2)$ time. The space impact is given by $O(n^2 \cdot |D|) = O(n^2)$ since there is at most one edge per direction per vertex.

Placing and removing a robot in the graph affects at most $2 \cdot n$ vertices. This is the case where the robot is placed where no obstacles or other robots exist for the given row and column. Finding the distant vertex uses $O(1)$ time. The total time for placing or removing a robot from the graph uses $O(2n) = O(n)$ time.

For the solver, each game state processed requires four robot placements and removals thus $O(8 \cdot n)$ time. The BFS worst case is given by the number of vertices and edges, resulting in $O(n^2 \cdot |D|)$. The BFS requires a two-dimensional array to keep track of the search progress using $O(n^2)$ space. However, this is discarded when the BFS is completed and the result is returned.

Validating, inserting, and enqueuing a new game state uses $O(1)$ time. Moving all robot uses $O(b)$ time. Therefore, the time used for each processed game state is $O(8 \cdot n + n^2 \cdot |D| + b) = O(n^2)$. In total, the worst case time usage for the algorithm is $O(b^k \cdot n^2)$.

The space impact of the queue and the hashtable is $O(b^k)$. The worst case space impact of the algorithm is $O(n^2 \cdot |D| + 2b^k + n^2) = O(b^k + n^2 \cdot (1 + |D|)) = O(b^k + n^2)$.

After the initial move, each OR is adjacent to at least one obstacle. The branching factor is given by $b = |OR| \cdot 3 = 9$.

2.3 Graph v2

The second version of the graph solver has a few improvements inspired by the other solvers. In the following, only the additions will be described.

2.3.1 Improvements

The graph construction, inserting, and removal of robots in the graph is the same as bescribed for version 1 while the same basic assumptions as stated in version 1 applies. The additions is focused on the heuristic function and the BFS while the solver itself has no further additions.

2.3.1.1 Heuristic Function

The heuristic function has been improved to facilitate an earlier termination of the solver. Using the minimum moves precomputation given in the IDDFS solver, the heuristic function has been updated to the following:

> B is the optimal solution if $B.best < (s.moves + min[GR])$, where s is the curent investigated game state and *min[GR]* is the minimum number of moves from the *GR* to the goal. Given the incremental order of the search, no possible other solution can reach goal in fewer moves than the given best, and the condition is satisfied.

2.3.1.2 Early Termination of BFS

When a given solution is found, the BFS is still applied until the heuristic function is satisfied. However, when a solution is found an upper limit for the possible number of moves is set drastically reducing the search tree. The following optimizations are an attempt to do so:

1. The depth of the BFS can be limited by the best solution so far and the minimum required moves to the goal. Let $depth = B.best - s.moves$, then it's given that $depth - 1$ is the maximum searchable depth in the BFS that can improve the current result. The BFS is terminated if this condition is fulfilled.
2. Let c be the current node processed in the BFS and let $c.d$ denote the distance to the root. If $(min[c] + c.d) \geq depth$ is true, the current searched branch in the BFS tree is obsolete and the rest of this branch is skipped. However, the BFS is not terminated.

2.3.2 Analysis

The minimum moves computation is the only addition to the worst case analysis. As already described in the IDDFS solver, it uses $O(n^3)$ time and $O(n^2)$ space. Therefore, the total worst case time usage is $O(n^3 + b^k \cdot n^2) = O(b^k \cdot n^2)$. Since the branching factor continues to be the dominant factor, the worst case time remains the same. The space impact of the algorithm is $O(b^k + n^2 \cdot (1 + |D|) + n^2) = O(b^k + n^2 \cdot (2 + |D|)) = O(b^k + n^2)$. The intermediate result shows a slight increase in the space impact while the worst case remains the same as version 1.

2.4 Summary

In table 1 all described algorithms are summarized with branching factor and worst case space and time analysis. As shown, only the stateless algorithm does not include the branching factor in the running time and space analysis.

Table 1: A summary of the presented algorithms.

Algorithm	Branching factor	Worst Case Time	Worst Case Space
Naive	~ 12	$O(b^{k+1} \cdot n)$	$O(b^k + n^2)$
IDDFS	~ 12	$O(b^k \cdot n + n^3)$	$O(b^k + n^2)$
Stateless	-	$O(n^3 \cdot log(n^2))$	$O(n^2)$
Graph v1	~ 9	$O(b^k \cdot n^2)$	$O(b^k + n^2)$
Graph v2	~ 9	$O(b^k \cdot n^2 + n^3)$	$O(b^k + n^2)$

The worst case running time of the naive algorithm is the worst while the state-less is the fastest in the worst case and most optimal in the space analysis. The graph algorithms are both much alike and only the minimum moves addition to version 2 affects the worst case time. The IDDFS algorithm has the best worst case running time of the four algorithms depending on the branching factor.

3 Experimental Results

3.1 Setup

I have compared the following algorithms.

- **Naive** The basic algorithm for the problem in own implementation.
- **IDDFS** The IDDFS algorithm claimed to be the best solver in own implementation.
- **Stateless** My own proposed stateless algorithm independent of the branching factor.
- **Graph v1** My own graph and dynamic programming based algorithm.
- **Graph v2** The optimized version of the one above. It applies many early termination techniques aggressively to shorten the average running time.

All performance tests have been conducted on a Intel i7 2.3 GHz with 4 GB heap allocated.

All of the described algorithms were implemented in Java version 1.8.0_60 and are available on GitHub[3]. All performance tests have been on the same game board but with different initial game states. The gameboard is available on appendix A. Given a game board size of 16^2 and 20 illegal starting fields for robots, there exist ~3 billion possible combinations.

In the following, the combination of a initial game state and a given goal is referred to as a game configuration. For each initial game state, 16 goals exist. Each game configuration is processed by all algorithms in a sequential manner. An upper time limit at 40 seconds is introduced to terminate long running

processes; at this point heap size becomes an issue and slows down the processing even further making the result ambigious. The result of the IDDFS is used as optimal baseline for each game configuration.

I have conducted performance tests on 160,000 game configurations. In the following, *required number of moves* will refer to the minimal number of moves required for reaching goal for the given robot while the *hitrate* will refer to the percentage of game configurations where a given algorithm solves the game configuration in the required number of moves. The hitrate and running time will be the primary focus of the tests.

3.2 Results

Using the result of the IDDFS a distribution of game configurations per required number of moves is shown in figure 7. Here it is seen that the the current board averages around 6-7 required moves per game configuration. It also shows few solutions exist which require 17 moves. Solutions having 17 required moves to goal are therefore pruned from the data sets given the insignificance and inaccurancy of these.


Figure 7: The distribution of game configurations.

As shown in figure 8, the naive algorithm reaches the 40 seconds limit at about 13 moves while the IDDFS solves every configuration. Also, the stateless and graph solutions follows each other closely with a 5-10% better hitrate for the latter. However, as the game configurations become more complex, the hitrate drops to ~50% for both solutions. Even though it seems that the hitrate is low, the graph algorithms have a total accuracy close to 90% while the stateless algorithm solves the game in ~80% of the game configurations as shown in figure 9.

In figure 10 the average running time is plotted as a function of the required number of moves for the given game configuration. The naive algorithm is

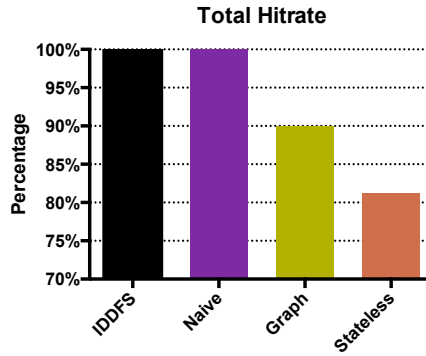Figure 8: The hitrate as a function of required number of moves.

Figure 9: The total hitrate as a percentage.

the slowest algorithm by several degrees and has been left out. The stateless algorithm uses 1-2 ms per solution indifferent to the number of required moves while the graph v2 algorithm is considerable faster than the IDDFS at nine moves required and higher.

Several results for the graph algorithms are not optimal and therefore the average running time given in figure 10 includes results where the algorithms investigated more moves than actually required. In figure 11 only the running time of results where the graph algorithms found the optimal solutions are compared. The stateless algorithm is omitted given the running time is not affected by the required number of moves.

All data is available on appendix B.

Figure 10: The average running time as a function of required number of moves.

Figure 11: The average running time as a function of number of moves given by the result.

3.3 Analysis

The running time of each algorithm corresponds well to the theoretical analysis. All algorithms depending on the branching factor uses a considerable amount of time as the complexity of the game configurations increases. For the stateless algorithm, the worst case running time is constant given the fixed size of the board, which also shows in the results.

For the given game board, the result of 80% of the game configurations can be found in constant worst case running time.

It is clear that the optimizations applied to the graph algorithm have affected the average running time consideredly. The extra cost of precomputing the minimum moves is not a considerable factor. Comparing the two graph algorithms, the increased performance of the graph v2 is an exponential function of the number of moves. This corresponds well to the actual optizmizations applied.

An interesting fact is that the graph v2 algorithm outperforms the IDDFS on solutions with 9 or higher required number of moves. For solutions where the graph v2 algorithm found the optimal solution, it outperforms IDDFS on game configurations with 8 or higher required number of moves. On solutions which required 16 moves, the graph v2 is faster with about a factor 7.5 if considering only the cases where the graph v2 found the optimal result.

However, the results also show that as the graph v2 and the stateless algorithms outperform the IDDFS algorithm on solutions with nine or higher required number of moves, the hitrate for each algorithm decreases fast. Therefore, in the search space where the proposed algorithms outperforms the known solutions, their reliability drops.

4 Future Work

Given the results of my work, I would like to propose several improvements and optimizations.

4.1 Performance Tests

The performance tests are conducted on the same game board. However, it would be interesting to investigate other game boards with a different distribution of game configurations per required number of moves and how it affects the running time of each proposed algorithm.

As the prototypes are implemented in Java, the algorithms could be implemented in C/C++ for increased performance while several machine-aware techniques could be used to optimize all algorithms in general.

4.2 Stateless

In few cases, the stateless algorithm returns a lower solution than actually possible. This means that there exists cases that are not accounted for, and these would have to be included in the assumptions. If these cases can be included without affecting the running time, the stateless algorithm will always guaranteee an upper limit for the required number of moves. If this is the case, it can be included in other algorithms and used to prune the search field.

4.3 Graph

During my tests I found that the BFS visits ~60 nodes for graph v1, while for graph v2 it only visits ~15 on average. This number directly depends on the number of reachable fields on the board which again depends upon the number of obstacles. A more precise worst case running time analysis could be proposed. This would depend on a function of the number of fields and the number of obstacles on the board instead of the number of nodes.

Another interesting improvement would be to move all robots including the *GR*. This would increase the branching factor, but the graph algorithm would guarantee an optimal solution. Given the results of the performance tests when comparing the average running time when the graph algorithms found the optimal solution, there is still a considerable gap between the IDDFS and the graph v2. A tradeoff between increasing the running time while guaranteeing the optimal solution could be of interest.

The pruning techniques introduced in the graph v2 could be improved. They are not truely applied before the first time the BFS hits goal. Thus, the graph

v2 relies on finding an upper limit before the pruning techniques works optimally. If the stateless algorithm can be improved to guarantee an upper limit for the required number of moves, this might be implemented into the graph algorithm when solving complex game configurations, setting an upper bound for the search field and initiate the pruning earlier.

4.4   IDDFS

The IDDFS could introduce the stateless algorithm on the more complex game configurations to improve overall performance and skipping some intermediate levels of deepening. However, given the incrementing nature of the algorithm, it does not benefit the same way as the graph v2 does. Depending on the inaccuracy of the stateless algorithm, this improvement could improve the running time on the more complex algorithms.

References

[1] Ricochet Robots: Solver Algorithms, https://speakerdeck.com/fogleman/ricochet-robots-solver-algorithms, 3/8 2015.

[2] Birgit Engels, Tom Kamphans: Randolphs Robot Game is NP-Hard! Electronic Notes in Discrete Mathematics, Volume 25, 1 August 2006.

[3] Github Source Code Repository, https://github.com/RydbirkJr/RicochetRobots

A   Game Board

Figure 12: The game board used for testing. The thick lines are obstacles. The colored fields are goals in the given color. The robots can be placed on all white fields.

B   Test Results

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
6958	5100	9920	15503	23860	30027	28370	19251	10579	5309	2699	1380	688	264	78	13	1

Table 2: The distribution of game configurations per required number of fields.

Algorithm	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
Naive	100	100	100	100	100	100	100	100	100	100	100	100	95.64	84.47	66.67	30.76
IDDFS	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100
Graph (v1 + v2)	100	100	100	98.37	93.40	90.06	88.63	85.51	80.82	74.51	67.25	57.10	53.05	53.03	60.26	38.46
Stateless	100	98.61	93.85	90.07	83.14	81.96	77.99	73.71	69.39	65.49	59.80	56.74	56.25	48.86	48.72	38.46

Table 3: The hitrate as a function of required number of moves. All numbers are in percentage.

Algorithm	IDDFS	Naive	Graph (v1 + v2)	Stateless
%	100	99.93	89.88	81.13

Table 4: The hitrate as a function of required number of moves, all numbers are in percentage.

Algorithm	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
IDDFS	0	0	0	0	0	1	4	11	28	63	155	414	998	2192	2996	7299
Graph v1	0	0	0	1	5	12	27	60	132	281	611	1163	1872	2525	3158	5195
Graph v2	0	0	0	0	1	2	5	11	23	50	114	229	407	620	657	1454
Stateless	1	1	1	1	1	1	1	1	1	1	1	1	2	2	3	1

Table 5: Average running times as a function of required moves. All numbers are in *ms*.

Algorithm	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
IDDFS	0	0	0	0	0	1	4	11	28	63	155	414	998	2192	2996	7299
Graph v1	0	0	0	0	3	7	19	41	82	157	305	583	1052	1717	2593	3724
Graph v2	0	0	0	0	0	1	3	6	14	28	55	107	207	362	593	909
Stateless	1	1	1	1	1	1	1	1	1	1	1	2	2	2	1	1

Table 6: Average running times as a function of required moves limited only to results where the optimal solution was found. All numbers are in *ms*.